

DTCC

2015中国数据库技术大会

DATABASE TECHNOLOGY CONFERENCE CHINA 2015

大数据技术探索和价值发现



HAWQ

MPP SQL for HDFS of Hadoop

基于Hadoop原生HDFS的大规模并行SQL

HAWQ Is The...

Enterprise platform that provides the fewest barriers, lowest risk, most cost effective and **fastest** way to enter in to big data analytics on HDFS of Hadoop



HAWQ 简述

ANSI SQL 2003/2011 Support



Greenplum database re-platformed on Hadoop/HDFS

SQL Engine

Cost-Based Query Optimization

Robust Query Optimizer

Complex Data Management

Sub-Partitioning

Distributions

Partitioning

CPU

Mem

Disk

Users

Storage Options

Polymorphic Storage

Row/Columnar Storage

Built-in Compression

HDFS Native Formats

MapReduce Integration

Multi-User Platform

Resource Queues

Concurrency

Data Encryption

Role-Based Security

Accessibility

ODBC/JDBC Driver L3,4

Parallel Loading/Unloading

HDFS Native Formats

Extendable...

txt

Avro

Seq

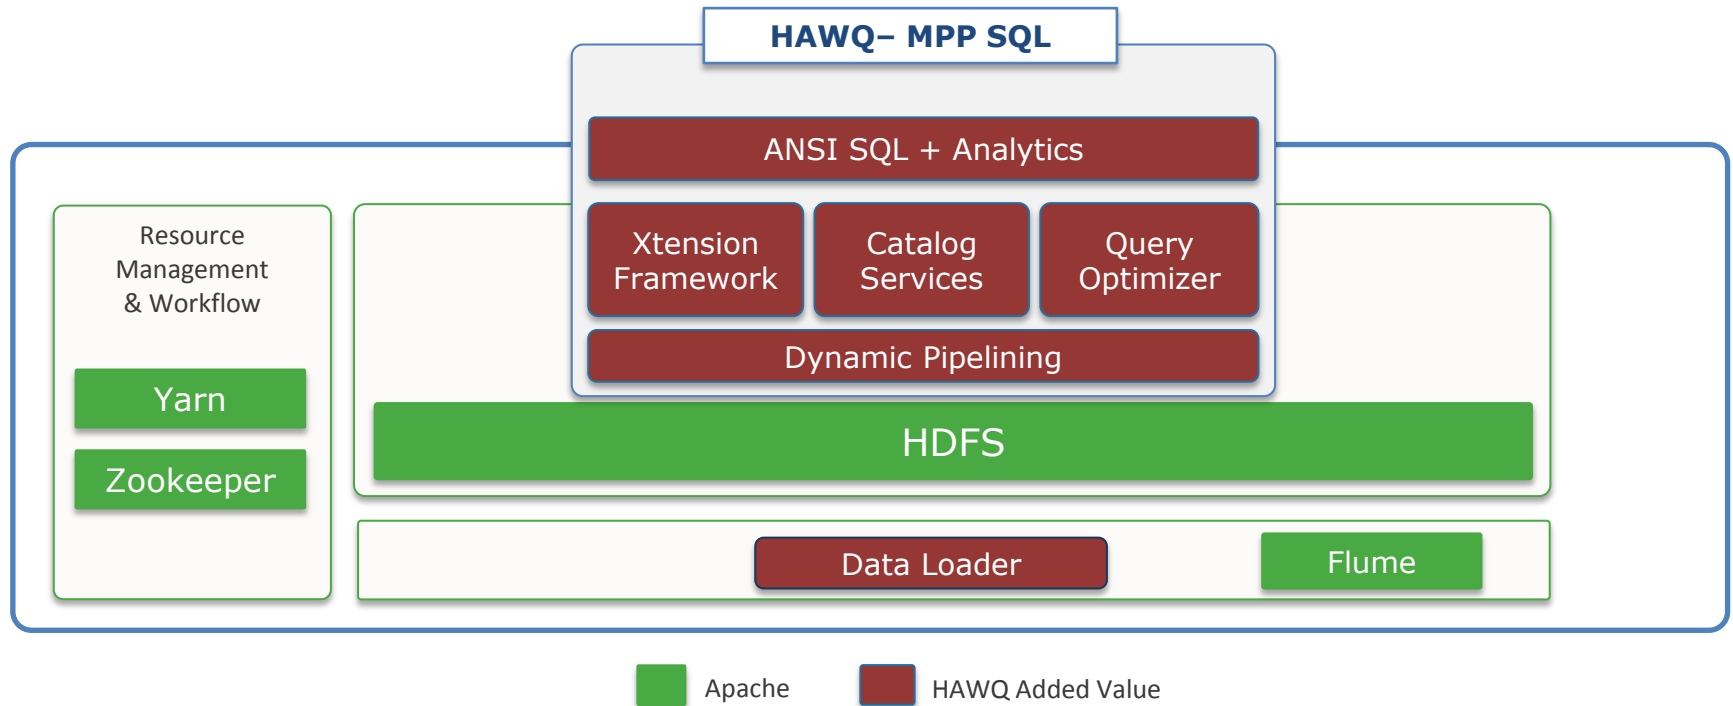
HBase

Hive


HAWQ的优点...

- 支持Apache Hadoop原生HDFS的SQL大规模并行引擎（MPP SQL）
- GPFX External Tables 接口，使用SQL透明访问Hadoop上各类数据
 - HDFS, HBase, Hive, Parquet格式等等
- 还支持SQL透明访问NFS，HTTP其他格式的数据（可自定义）
- Performance and Scalability
 - Parallel Everything
 - Dynamic Pipelining
 - High Speed Interconnect（基于UDP）
 - HDFS access with C++ libhdfs3
 - Co-Located Joins & Data Locality
 - Partition Elimination（支持静态动态表分区）
 - Higher Cluster Utilization
 - Concurrency Control（资源作业优先级调度）

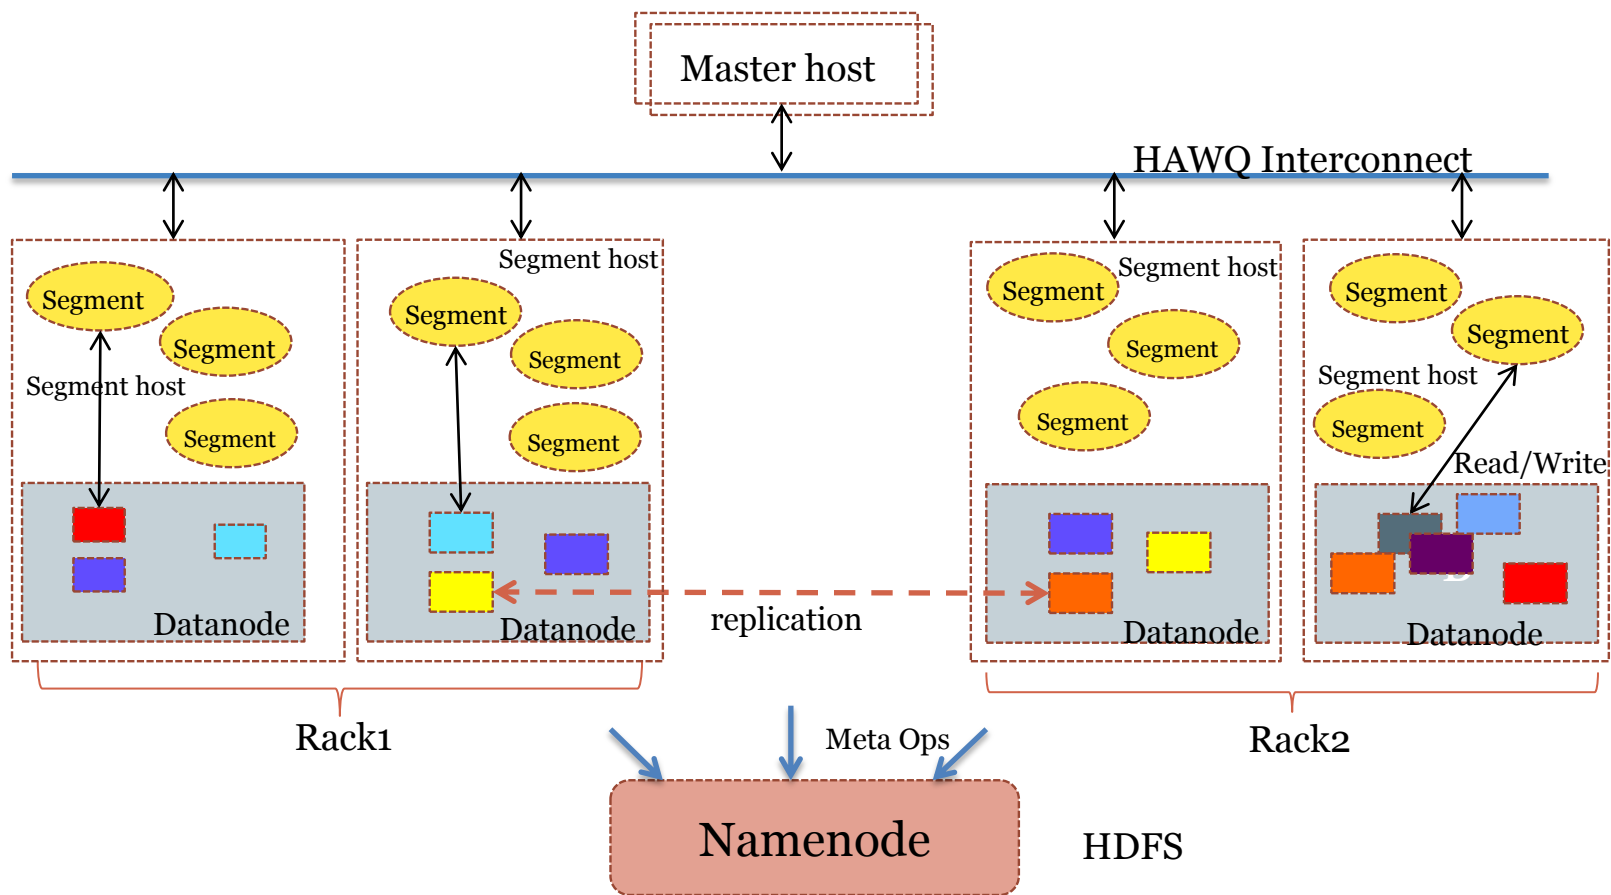
HAWQ 及Hadoop软件栈



HAWQ 与 Hadoop HDFS



HAWQ 与 Hadoop HDFS数据访问流

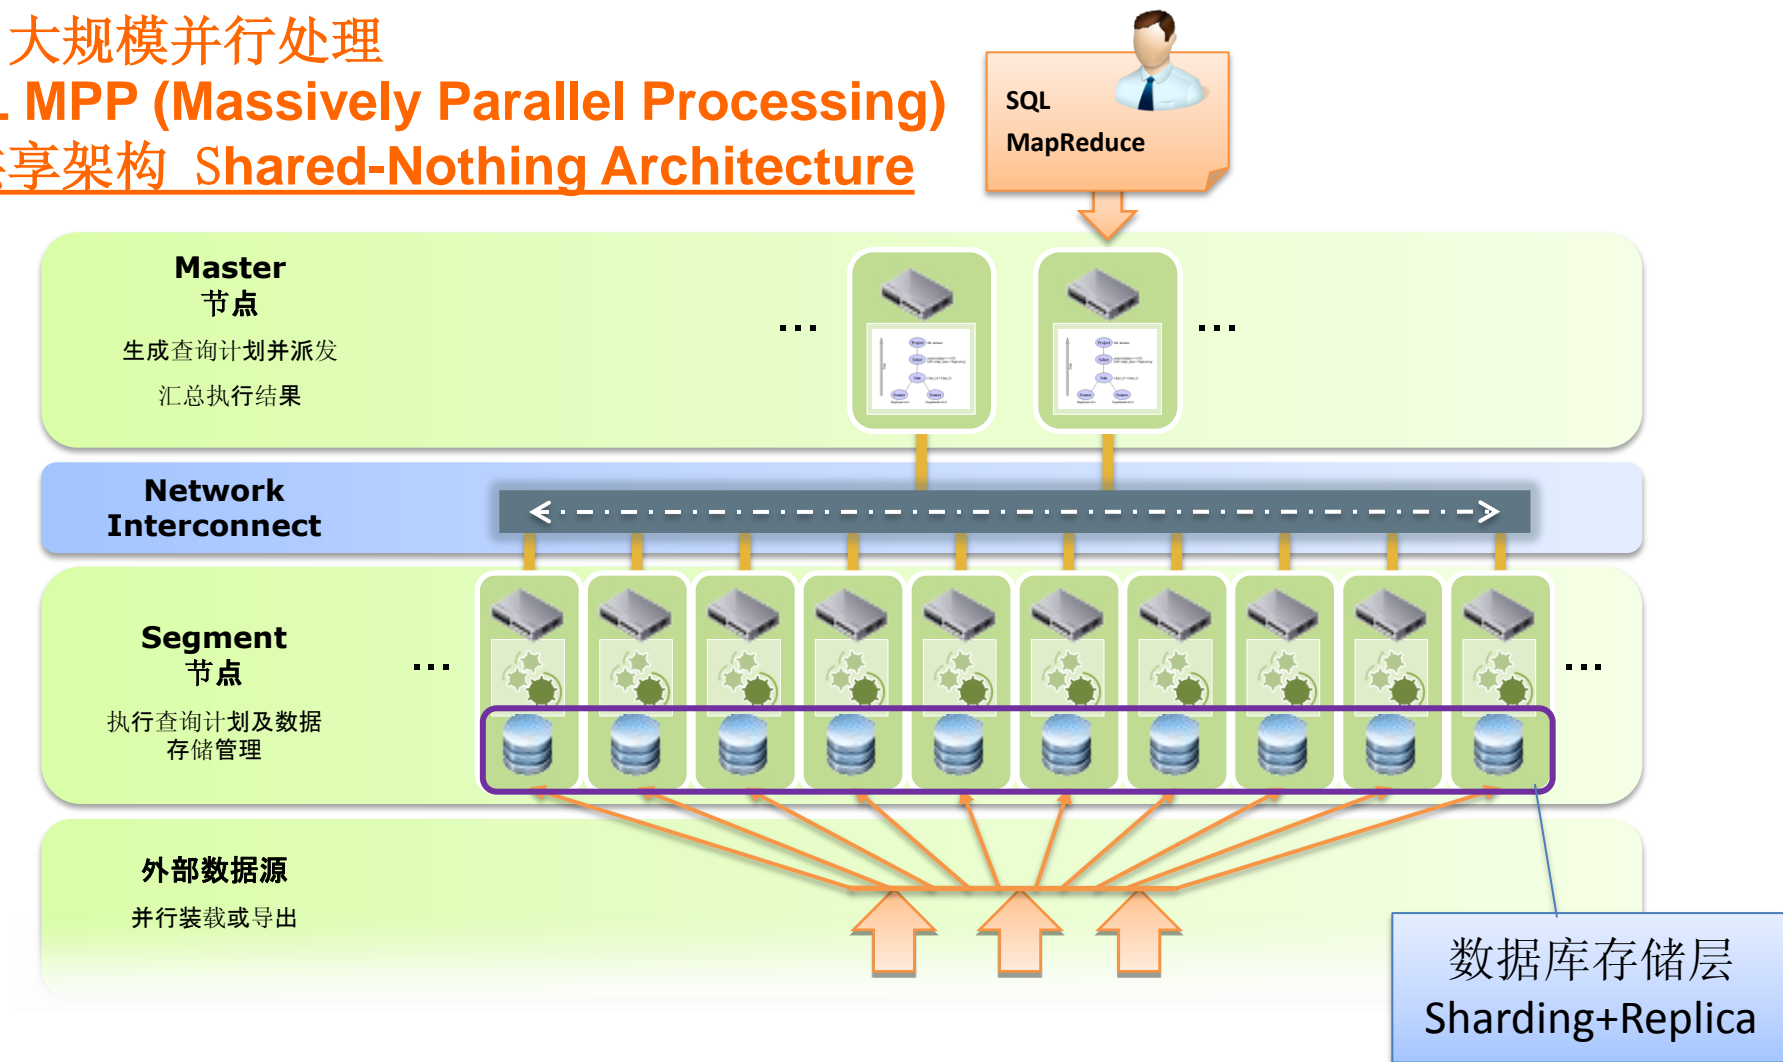


HAWQ 对比 Greenplum DB 基本架构

SQL 大规模并行处理

SQL MPP (Massively Parallel Processing)

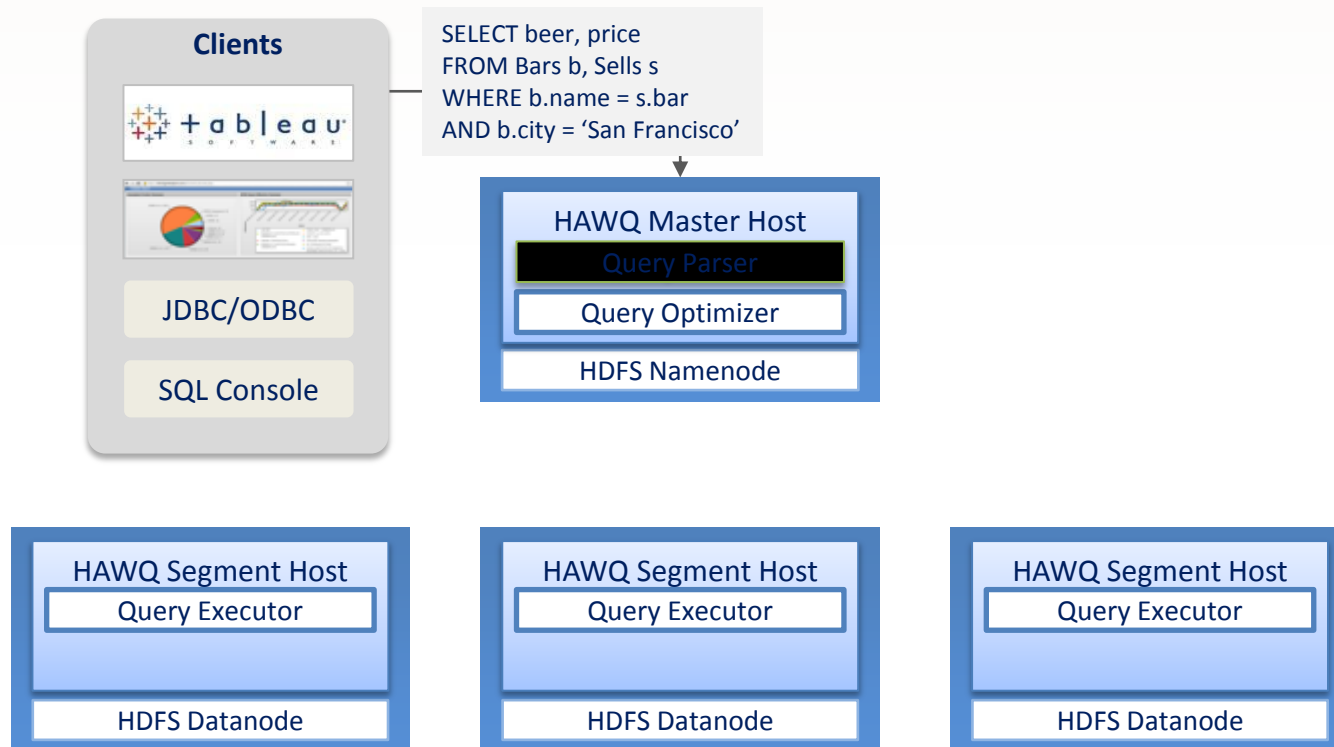
无共享架构 Shared-Nothing Architecture



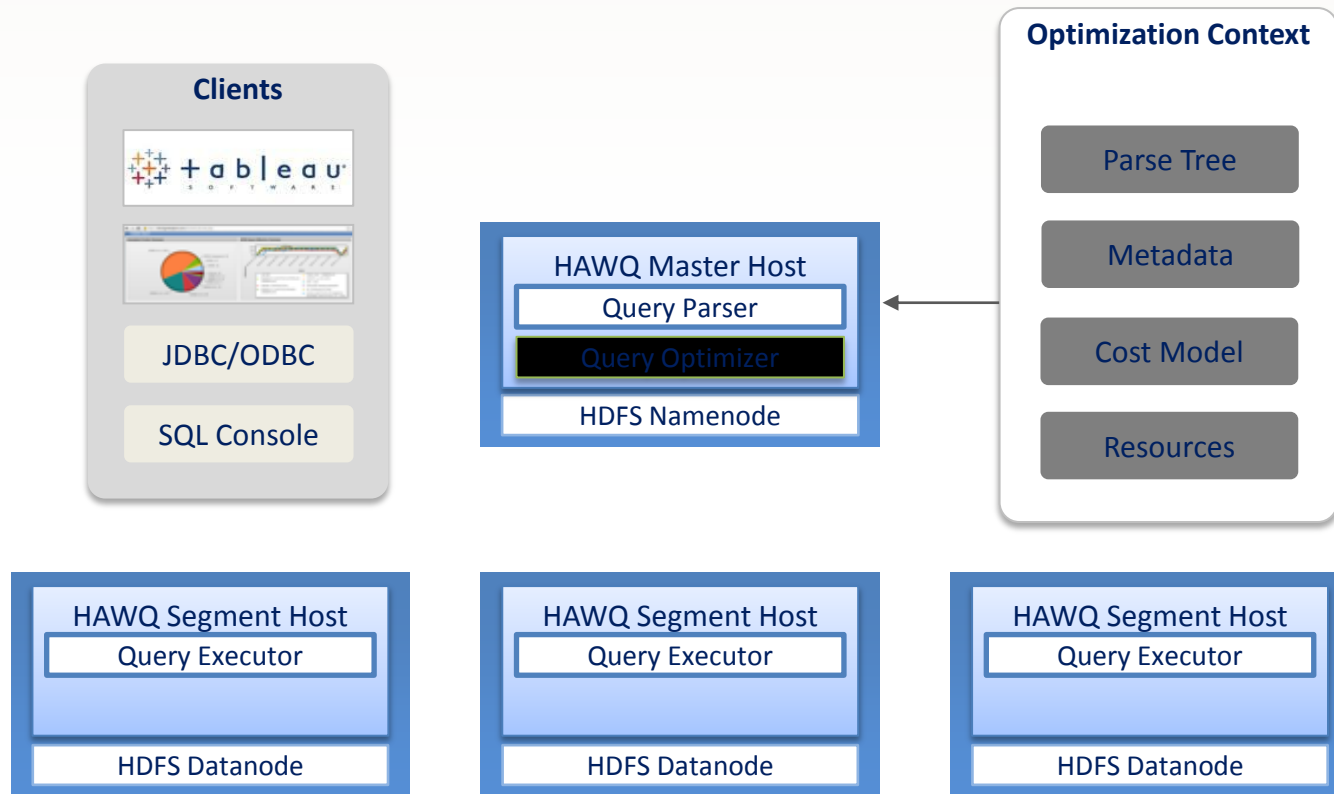
运行SQL，支持SQL2008及OLAP选项

```
-- Top 10 products for each category for all time
SELECT lineitems.product_id
,       products_dim.product_name
,       product_category_id
,       category_name
,       product_count
,       category_rank
FROM (SELECT product_id, product_category_id
,          SUM(item_quantity) AS product_count
,          row_number() OVER (PARTITION BY product_category_id ORDER BY SUM(item_quantity) DESC) AS category_rank
FROM   retail_demo.order_lineitems
GROUP BY product_id, product_category_id
) AS lineitems |
INNER JOIN retail_demo.categories_dim on product_category_id = category_id
INNER JOIN retail_demo.products_dim on lineitems.product_id = products_dim.product_id
WHERE category_rank <= 10
ORDER BY product_category_id, category_rank
limit 100;
```

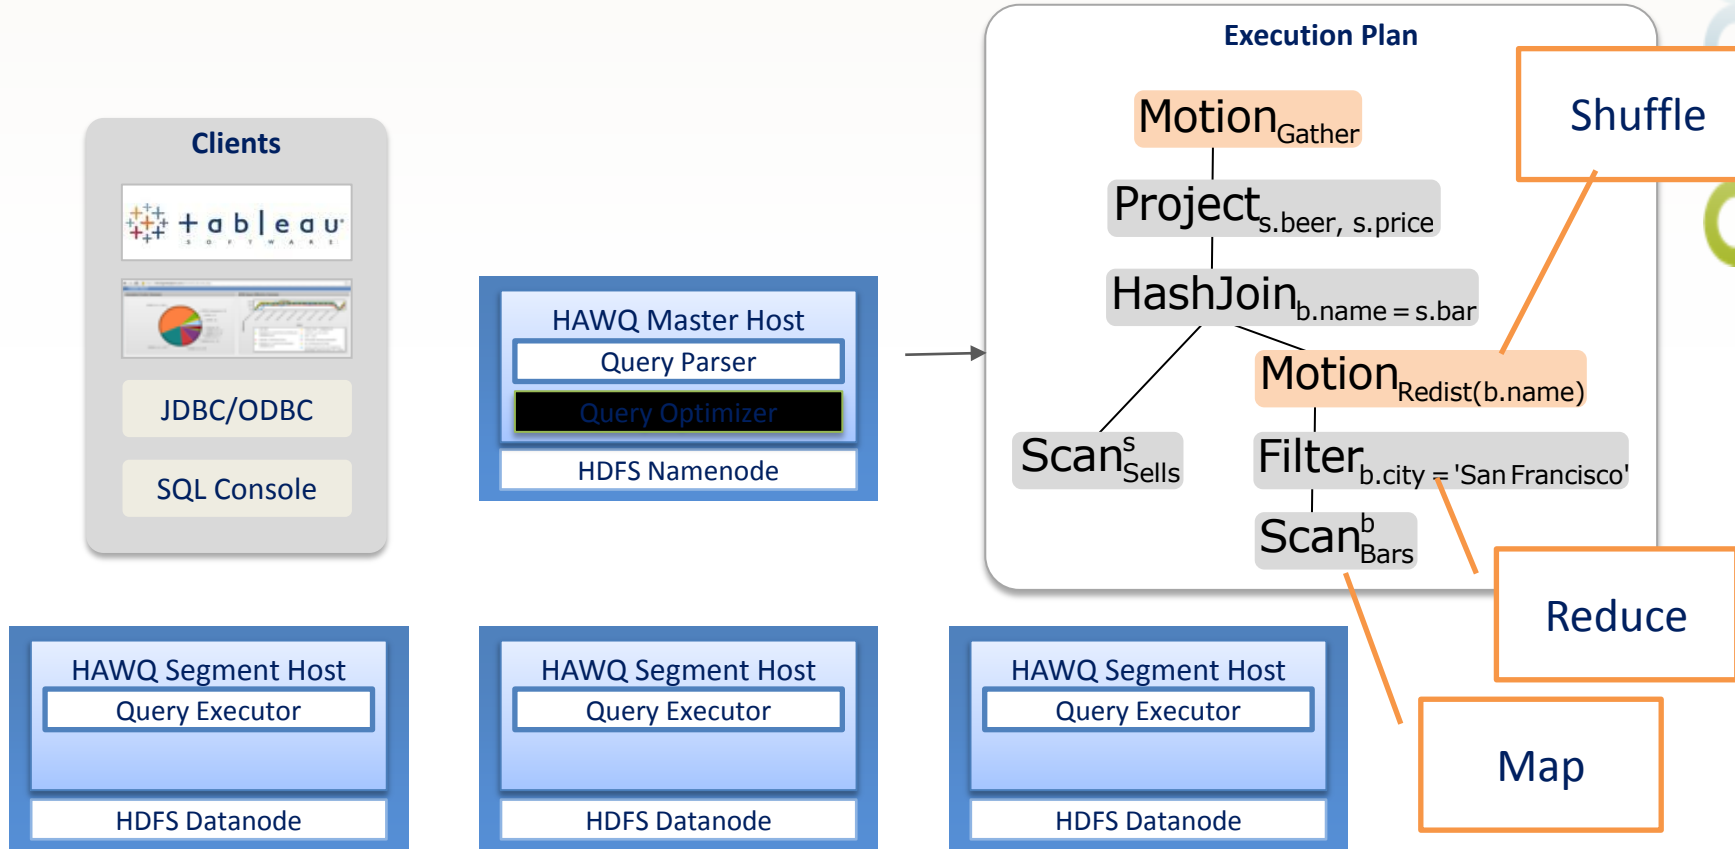
HAWQ (SQL MPP) 机制-1



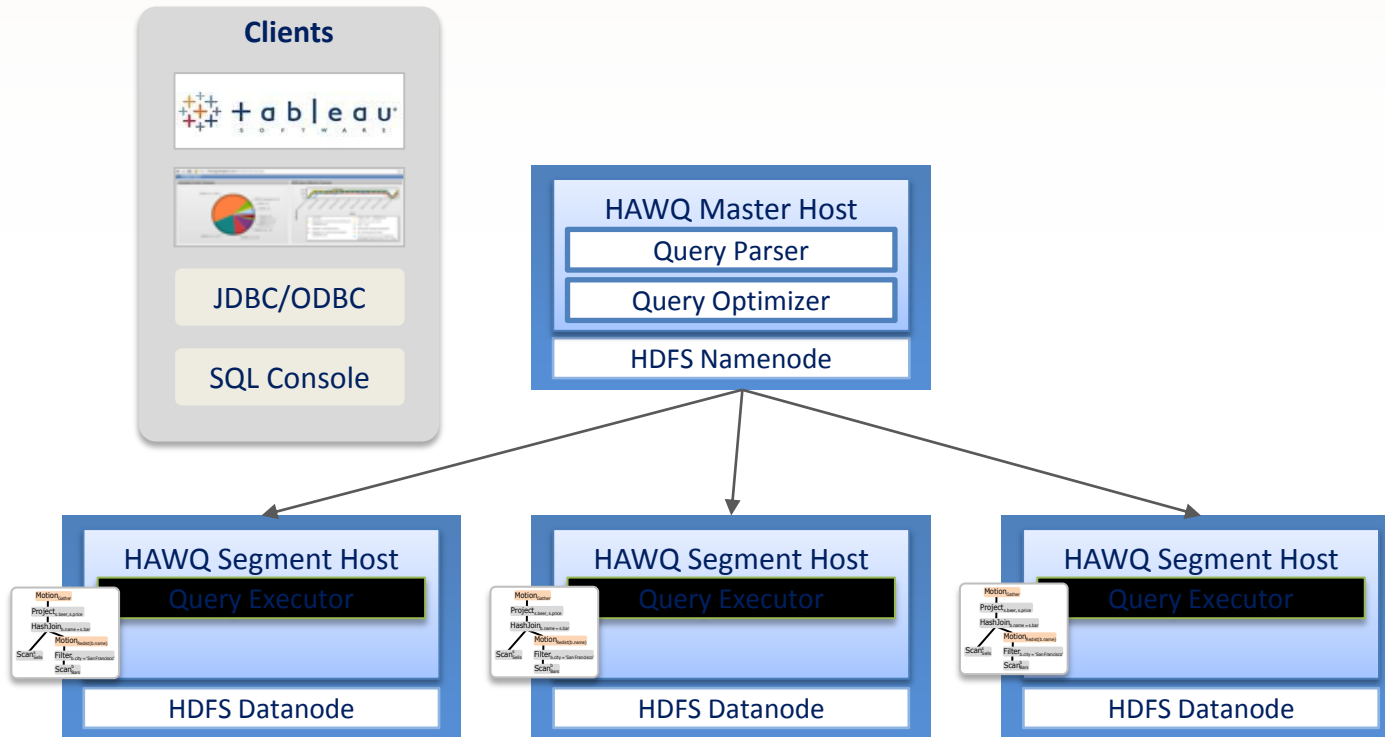
HAWQ (SQL MPP) 机制-2



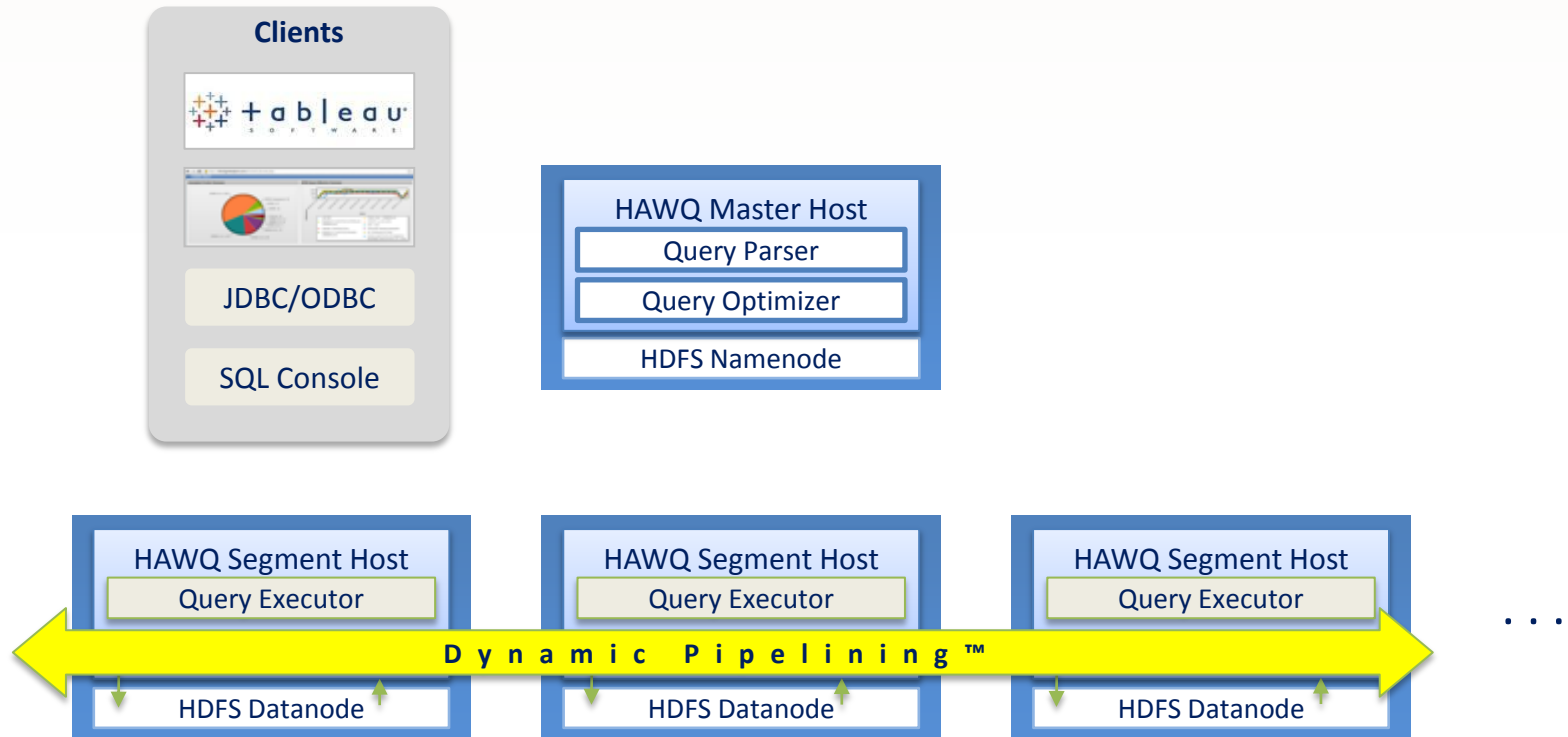
HAWQ (SQL MPP) 机制-3



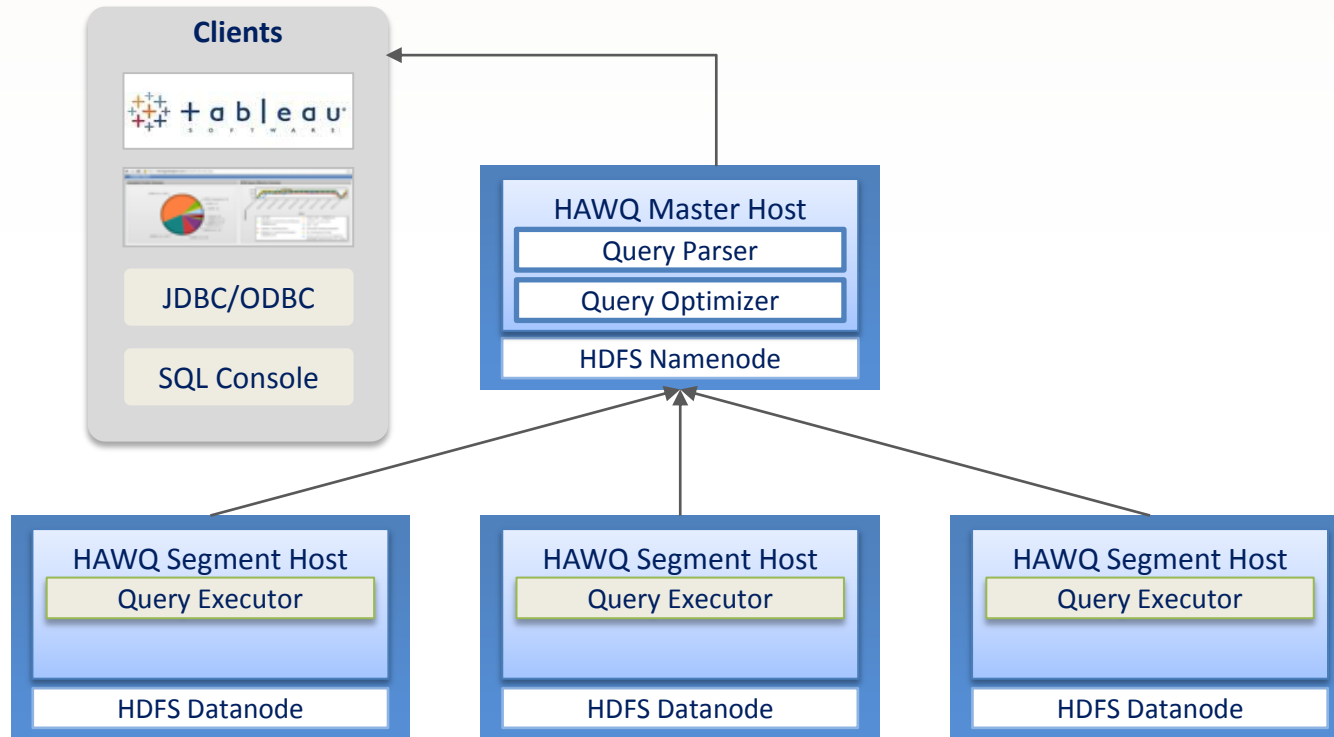
HAWQ (SQL MPP) 机制-4



HAWQ (SQL MPP) 机制-5



HAWQ (SQL MPP) 机制-6




数据分布方式(Data Distribution)

- Data can be distributed based on a column or a composite of columns
- Tables distributed similarly are **co-located**
- Distribution scheme modifiable thru alter table

Advantages:

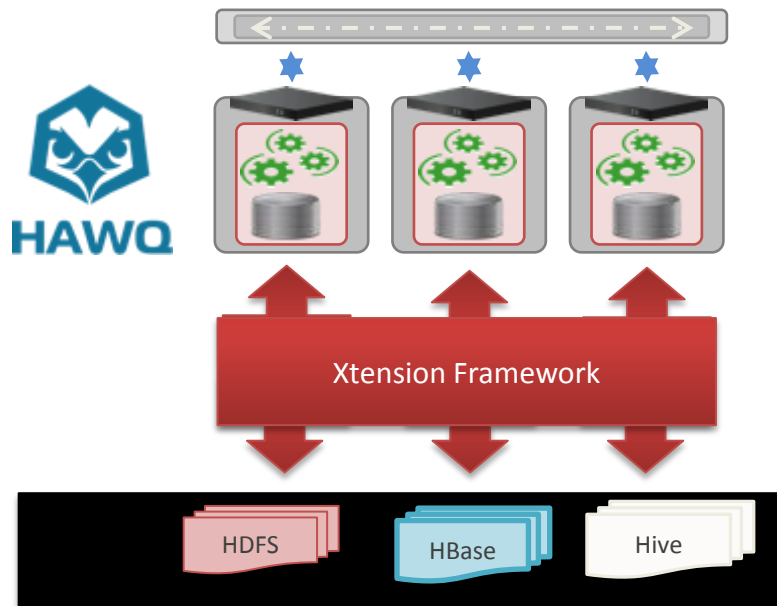
- Co-located joins
- No data movement on joins or aggregates
- Improved performance on complex queries
- Query engine optimization



```
SELECT X FROM A,B WHERE A.X = B.Y
```

```
SELECT SUM(X) FROM A GROUP BY A.X
```

数据互联框架(Xtension Framework)



- An advanced version of Greenplum DB external tables
- Enables combining HAWQ data and Hadoop data in single query
- Supports connectors for HDFS, Hbase and Hive
- Provides extensible framework API to enable custom connector development for other data sources

数据导入导出(Loading/Unloading Data)

gpload, gpfdist, External Tables 

Flat Files, CSV, Delimited, ...

Existing RDBMS Systems 

Web Tables, JSON, XML, HTML, ...

Executing Scripts, ... 

DataLoader 

File Farms 

Streaming || Batch Mode 

Flume, ... integration

Throttling, Compression, ...
features

PXF {Native Hadoop Files} 

HDFS Flat Files, CSV, Delimited, ... 

Hive

HBase {w. predicate push-down} 

Avro, RCFile, SeqFile

Open extendable API 

Available on Github: Accumulo,
JSON,...

Spring XD 


Java Development Framework

Traditional Tools 

Postgres insert, copy, ...

ODBC + JDBC drivers 

Pivotal Data Dispatch {PDD}

Integration with ETL tools...

HAWQ 里数据导入导出仍是全并行

HAWQ External Tables



gpload, gpfdist, External Tables

Flat Files, CSV, Delimited, ...

Existing RDBMS Systems



Web Tables, JSON, XML, HTML, ...

Executing Scripts, ...



HAWQ and Hadoop Native File Formats



PXF {Pivotal eXtension Framework}



HDFS Flat Files, CSV, Delimited, ...



Hive

HBase {predicate push-down}



Avro, RCFile, SeqFile

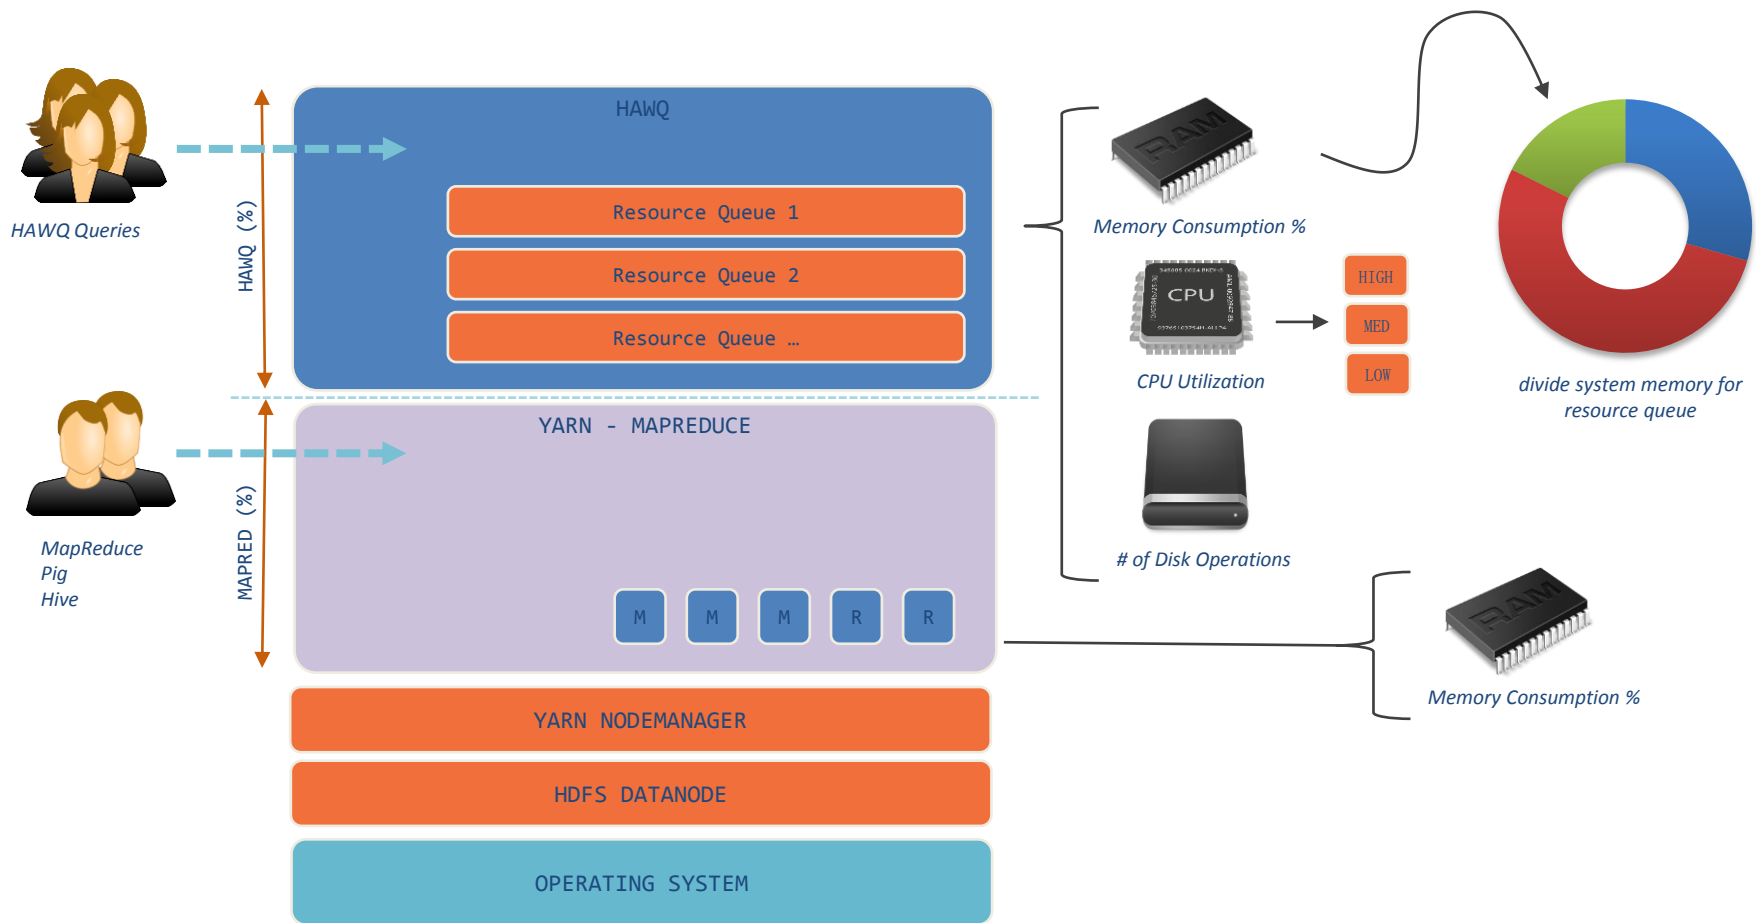
Open extendable API




Available on Github: Accumulo, JSON,...

Read/Write

更强大的资源管理器,兼容YARN



运行时资源可控 (Dynamic Resources Allocation)



SQL for Hadoop功能对比

Feature	Hive	Impala	HAWQ
Work with HDFS native file formats	✓	✓	✓
Polymorphic Storage	✓	✓	✓
Advance SQL (ANSI SQL2008 & OLAP support)	✗	✗	✓
Partitions and compression	✓	✓	✓
Data Locality	✓	✓	✓
Distributions, Join, Aggregate Locality	✗	✗	✓
Join Optimization	✗	✗	✓
Spill to disk (query must fit in memory)	✓	✗	✓
Fault tolerance during large query execution	✓	✗	✗
Granular Security and authentication	✗	✓	✓
Extendable (Serdes)	✓	✗	✓
Resource Management	✗	✗	✓
Open-source code	✓	✓	✓

HAWQ 性能对比-1



User intelligence	4.2	37	9X
Sales analysis	8.7	596	69X
Click analysis	2.0	50	25X
Data exploration	2.7	55	20X
BI drill down	2.8	59	21X

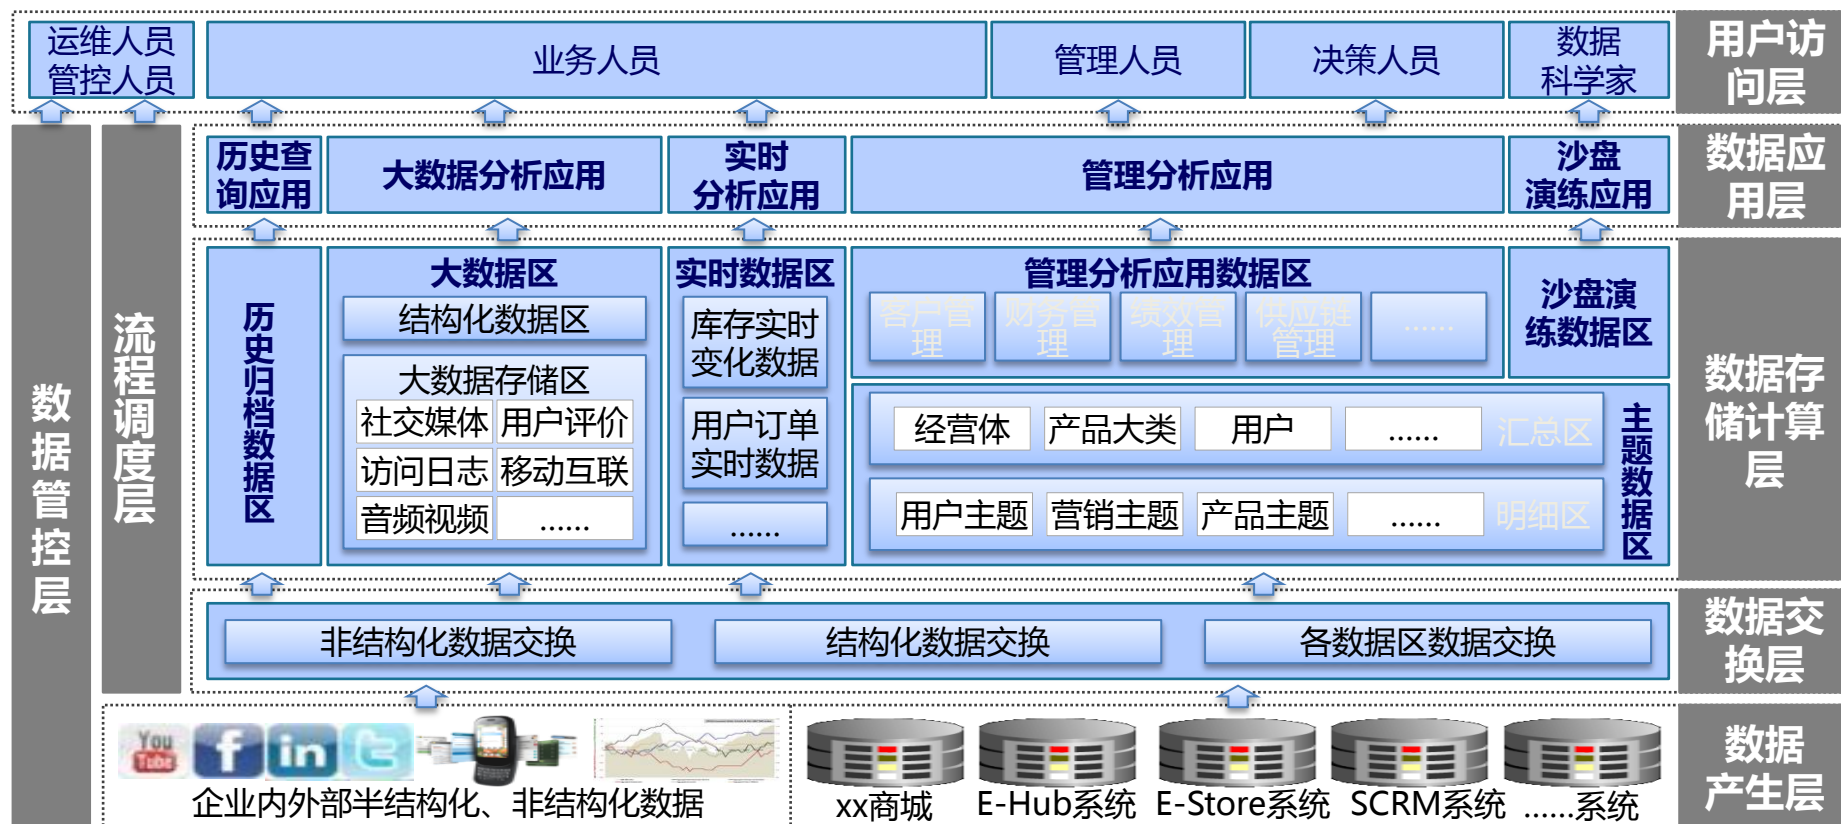
HAW性能对比-2



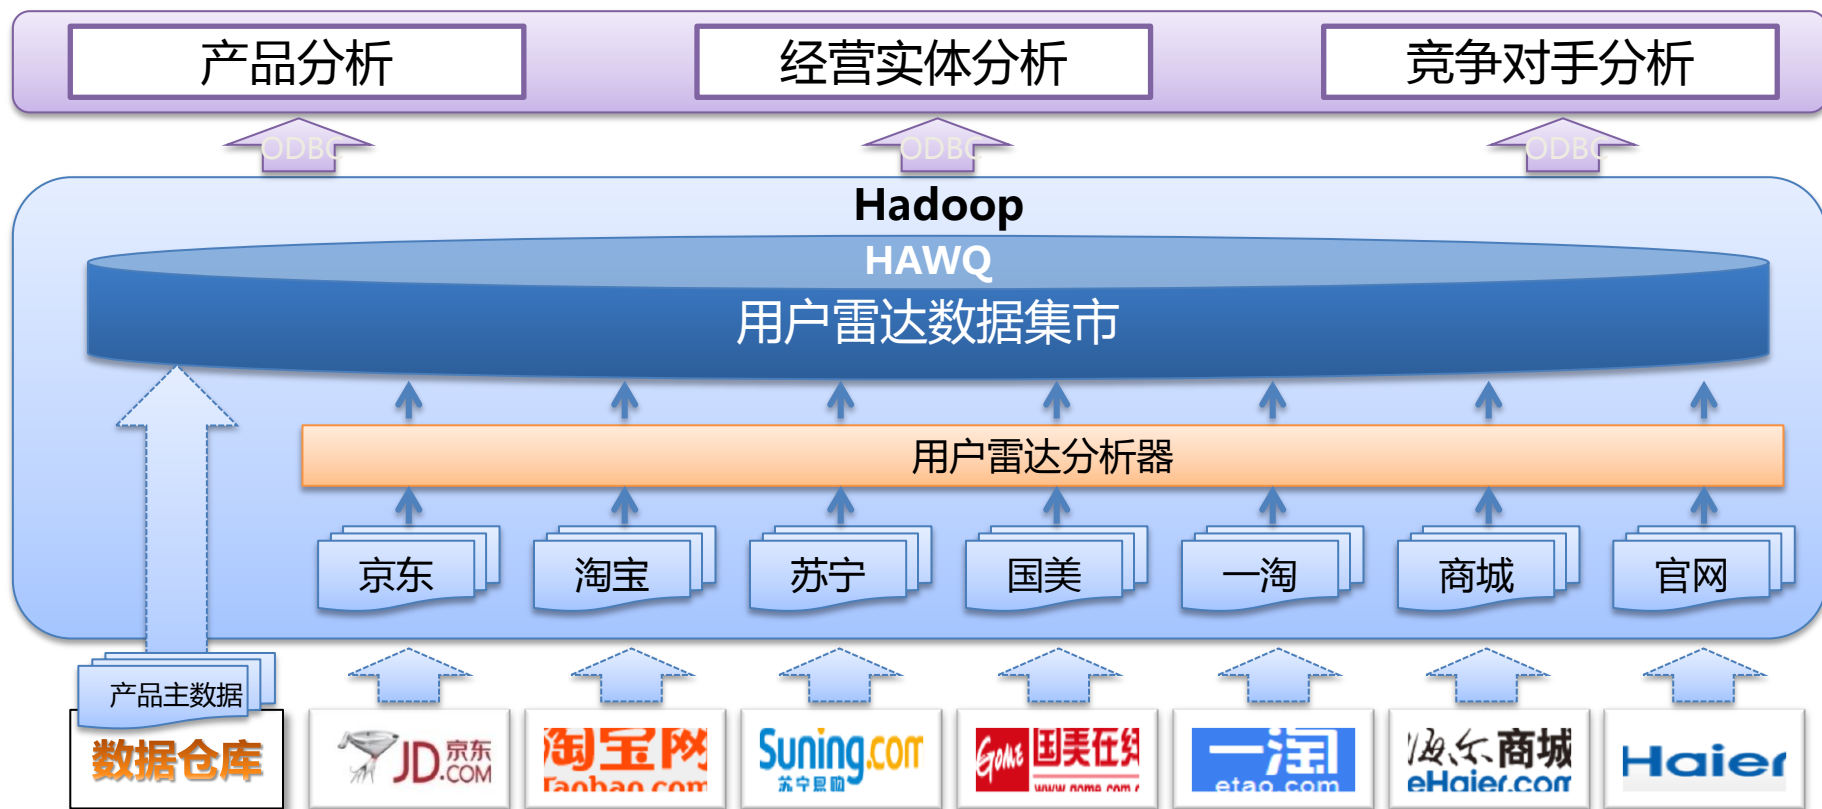
User intelligence	4.2	198	47X
Sales analysis	8.7	161	19X
Click analysis	2.0	415	208X
Data exploration	2.7	1,285	476X
BI drill down	2.8	1,815	648X

部分应用案例

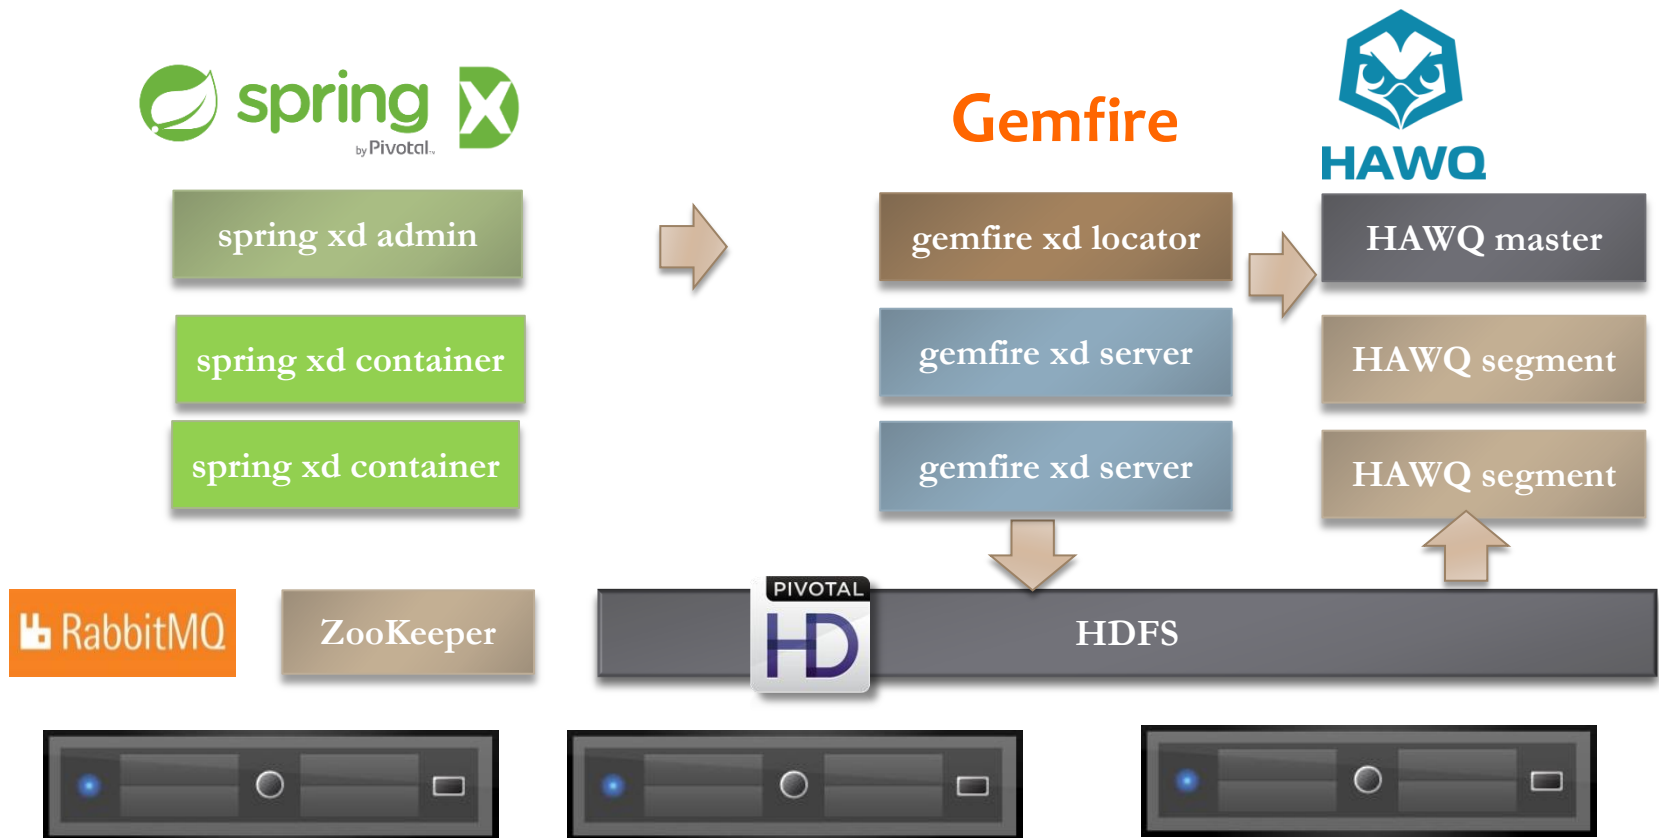
某企业解决方案



大数据分析应用：用户雷达



方案架构图





THANKS